

32-bit MIPS ISA



Introduction

- Four backward-compatible families: MIPS I, II, III, IV
 - We will focus on MIPS I only (32-bit ISA); what is this?
 - MIPS III onwards are fully 64-bit ISAs (in textbook)
- Load/store register ISA
 - Operates only on registers
 - Can access memory only through load/store (big endian)
 - No partial register access
- RISC philosophy
 - Emphasis on efficient implementation: Make the common case fast
 - Simplicity: provide primitives, not the solutions
 - A system can be so simple that it obviously has no bugs, or so complex that it has no obvious bugs [C. A. R. Hoare]

Data types

- Bit strings
 - Byte is 8 bits
 - Half word is 16 bits
 - Word is 32 bits
 - Double word is 64 bits
- Integers in 2's complement
- Floating-point: single and double precision
 - IEEE 754 standard

MAINAK CS422

1

Storage model

- 32-bit address: 4 GB addressable memory
- Separate 31x32-bit GPRs (\$0 is hardwired 0) for integer and 32x32-bit GPRs for floating-point
 - Writing to integer \$0 will not change it (these are NOPs)
- Program Counter (PC) is incremented by 4 (except branch, jump)
 - Instructions are 32-bit in size (for all four families)
- Two special registers Hi and Lo for storing multiply/divide results
- Floating-point registers are paired for doing double-precision
 - The pair \$f_{2n} and \$f_{2n+1} are accessed by name \$f_{2n} e.g. \$f2 specifies the 64 bits in \$f2 and \$f3 with the least significant word in \$f2

1

single: $\text{add.s } \$f_0, \$f_2, \$f_4$:

f_2
$+f_4$
f_0

double: $\text{add.d } \$f_0, \$f_2, \$f_4$

must be even for
(.d) commands

f_3	f_2
$+f_5$	f_4
f_1	f_0
1	

Computation

- ALU instructions
 - Classic 3-operand format: two sources and one dest.
 - Operands: GPRs or 16-bit immediate values
 - Both signed and unsigned arithmetic
 - Basic difference between signed and unsigned arithmetic: overflow not flagged for unsigned
 - In both cases, the operands are treated as signed
 - Sign extension of immediate in both signed and unsigned arithmetic; zero extension for immediates in logical inst.
 - Signed comparison is completely different from unsigned comparison (unsigned comparison treats the operands as unsigned; not true for unsigned arithmetic)
 - Integer multiply/divide take only two operands i.e. two sources and have implicit targets Hi and Lo: ISA offers instructions to move from Hi/Lo registers to GPR

MAINAK CS422

1

ALU instructions

- | | |
|---------------------------------|--------------------|
| • Arithmetic | • Logical |
| add \$3, \$2, \$1 | and \$3, \$2, \$1 |
| sub \$3, \$2, \$1 | or \$3, \$2, \$1 |
| addi \$3, \$2, 100 | xor \$3, \$2, \$1 |
| addu \$3, \$2, \$1 | nor \$3, \$2, \$1 |
| subu \$3, \$2, \$1 | andi \$3, \$2, 10 |
| addiu \$3, \$2, 100 | ori \$3, \$2, 10 |
| slt \$3, \$2, \$1 | xori \$3, \$2, 10 |
| slti \$3, \$2, 100 | sll \$3, \$2, 10 |
| sltu \$3, \$2, \$1 | srl \$3, \$2, 10 |
| sltiu \$3, \$2, 100 | sra \$3, \$2, 10 |
| mult/multu \$3, \$2 | slvl \$3, \$2, \$1 |
| div/divu \$3, \$2 // Lo=q, Hi=r | srlv \$3, \$2, \$1 |
| mfhi \$4 | sra \$3, \$2, \$1 |
| mflo \$4 | lui \$3, 40 |

1

Floating point


- Operates on floating-point registers
 - Supports both single and double precisions
 - No hardwired zero register
 - IEEE 754 compliant
 - Typical instructions
 - add, sub, mul, div, mov (one fp reg to another), neg, abs, cvt (precision conversion), mfc, mtc (move between fp and integer registers)
 - mul/div don't use Hi/Lo; target FPR is specified explicitly (remainder doesn't have a meaning here)

MAINAK CS422

1

Load/store

- One address mode for memory
 - Displacement only (always sign-extended)
 - Most loads/stores are aligned (except lwl/lwr, swl/swr)
 - Loads are supported for signed and unsigned data; unsigned loads zero-extend the loaded value
 - Supports three sizes: byte, half word, word (double word is supported in 64-bit ISA)
 - Byte load: lb \$3, 4(\$13) or lbu \$3, 27(\$20)
 - Half word load: lh \$3, 12(\$10) or lhu \$6, 0(\$7)
 - Word load: lw \$20, 52(\$29) [no unsigned flavor]
 - Byte store: sb \$3, 5(\$2)
 - Half word store: sh \$3, 56(\$29)
 - Word store: sw \$2, 20(\$2)
 - The immediate is *not* shifted by load/store size; it is just sign-extended and added to the base register content
 - In addition, floating-point load/store: lf \$f1, 60(\$22)

$\$10 = 57$ 

Direct addressing

- In some cases the address is known at compile-time
 - Happens mostly for statically allocated global variables
 - How do you emulate direct addressing?
 - Suppose I want to load from address 0x123456 (addr)

lui \$2, 0x12

ori \$2, \$2, 0x3456

lw \$4, 0(\$2)

Could save one instruction by using non-zero displacement

lui \$2, 0x12

lw \$4, 0x3456(\$2)

What if the address is 0x789abc?

MAINAK CS422


78 0000
+ ffff9abc
≠ 789abc

lwl/lwr

- Load-word-left and load-word-right
 - Example: lwl \$4, 0(\$10) // Suppose \$10 has 0x57
 - Let the word containing this byte address be w
 - Extract the bytes contained by w that *start* from this address (remember this is big endian)
 - Put these bytes (in this case one byte) in the *upper* portion of \$4 and leave the remaining bytes (in this case 3 lower bytes) unchanged
 - Example: lwr \$4, 0(\$10) // Suppose \$10 has 0x5a
 - Extract the bytes contained by w that *end* at this address
 - Put these bytes (in this case three bytes) in the *lower* portion of \$4 and leave the remaining bytes (in this case the upper byte) unchanged
 - Why are these useful?

MAINAK CS422

1



← 26 bits →
op target
26 bits ⇒ 2^{26} instr ⇒ 2^{28} B jump + (PC+4)

Control transfer

- Jump
 - Unconditional jumps and procedure calls use absolute address
 - MIPS ISA offers 26 bits to encode the absolute target; shift this target to left by 2 bits (4-byte instruction) and borrow the upper four bits of the next PC (i.e. PC+4) to form the complete 32-bit target $((PC+4) \& 0xf0000000) | \text{tgt} \ll 2$
 - The procedure call instruction (known as jump and link) saves the return PC (PC+8) in a fixed GPR (\$31)
 - Indirect jumps where target is not known (e.g. procedure return or case/switch or procedure call via function pointer) use jump register (jr) or jump and link register (jalr) instructions; both take a register operand where the target is found
 - Example: procedure return: jr \$31

MAINAK CS422

1


not PC but PC+4
PC+8 saved in \$31

Control transfer

- Conditional branches
 - All conditional branches are compare-and-jump type
 - PC-relative immediate offsets are always sign-extended
 - Examples:
 - beq \$1, \$2, 100 /* target = PC + 4 + sext(100 << 2) */
 - bne \$1, \$2, 100
 - bgez \$1, 100
 - blez \$1, 100
 - bltz \$1, 100
 - bgtz \$1, 100
 - Can use slt, slti, sltiu first followed by beq or bne with \$0
- All branch and jump instructions have a delay slot
 - Instruction right after the jump or branch is always executed

MAINAK CS422

1



PC+8 : because the j, jal, jalr, jr etc all come into effect one cycle later. So (PC+4) has already been executed. So we come back to PC+8. That PC+4 is always a **Branch Delay slot**

PC+4: because instr after PC has already in execution. So we find dist from PC+4.

1/30/2025

instr after branch will always execute before branch takes effect

Control transfer

```

count = 0;
for (i=head; i<=n; i++) {
    if (A[i]==key) count++;
}
    
```

Assume following allocation
 count in \$5
 i in \$6
 head in \$1
 start addr of A in \$2
 key in \$3
 n in \$4
 A is an integer array

```

addiu $5, $0, 0
addu $6, $1, $0
loop: slt $9, $4, $6
      bne $9, $0, exit
      → sll $7, $6, 2
      addu $7, $7, $2
      lw $8, 0($7)
      bne $8, $3, next
      → addiu $6, $6, 1
      addiu $5, $5, 1
next: j loop
      → nop
exit: ...
    
```

will always execute as after branch. So we blo i++ here. ct++ may not be done.

could not fill anything so nop.

MAINAK CS422 1

Procedure call

- MIPS calling convention
 - Parameters are passed in registers; extra parameters (more than 4) are spilled on stack
 - Compiler tries best to allocate local scalar variables in registers; stack is used for spilling
 - Two dedicated GPRs point to top of stack (stack pointer) and start of procedure frame (frame pointer)
- Global pointer
 - Global static scalar variables are allocated in a 64 KB static area at compile time
 - Accessed with \$gp + offset (these are determined at link time)

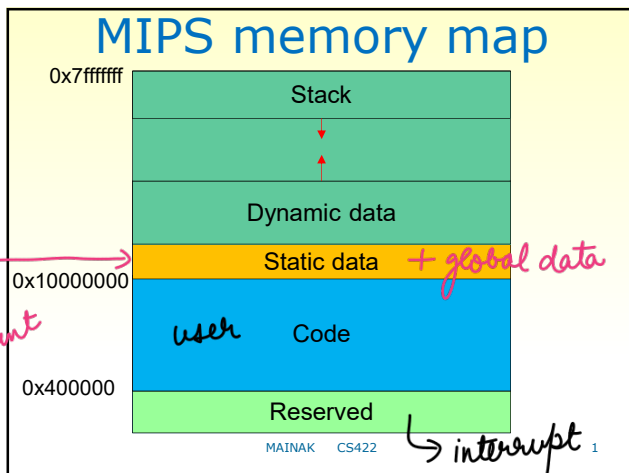
High order

\$sp → ↑ pop / ↓ push

non array

MAINAK CS422 1

Stack grows downwards



MIPS memory map

- User gets 2 GB
 - Upper 2 GB is kernel space
 - On processor reset execution starts at 0xbfc00000 (kernel space)
 - Stack grows downward while heap grows upward

MAINAK CS422 1

$\text{lui} + \text{ori} = \text{la} = \text{li}$ (needs temp reg)

Register convention

- 32 integer registers
 - \$0 is hardwired to 0 (really not implemented as register)
 - \$1 or \$at is reserved for assembler (la to lui conversion)
 - \$2, \$3 (or \$v0, \$v1) are return values of function with upper 32 bits in \$v1 and lower 32 bits in \$v0; \$v1 also holds the syscall number before executing system call
 - \$4 to \$7 (or \$a0 to \$a3) are procedure arguments; saved by caller
 - \$8 to \$15 (or \$t0 to \$t7) are temporaries; saved by caller
 - \$16 to \$23 (or \$s0 to \$s7) are callee saved
 - \$24 and \$25 (or \$t8, \$t9) are two more temporaries
 - \$26 and \$27 (or \$k0, \$k1) are reserved for kernel
 - \$28: global pointer (\$gp), \$29: stack pointer (\$sp), \$30: frame pointer (\$fp), \$31: return address (\$ra)

MAINAK CS422

1

Register saving

- Caller saves the registers that are not preserved across call
 - Of course, needed only if caller wants some value to be preserved *ret val*
 - arg* - \$a0 to \$a3, \$v0, \$v1, \$t0 to \$t9 *not reg to be saved*
- Callee saves the registers that are required to be preserved across call
 - Needed only if callee uses these registers
 - \$s0 to \$s7, \$gp, \$sp, \$fp, \$ra
- MIPS gcc cross-compiler and native cc compiler do not use frame pointer
 - \$30 is treated as callee-saved \$s8

MAINAK CS422

1

$\$1 = \$at = \$\text{ assembler temporary.}$
 $\$26, \$27 \rightarrow \text{return from exceptions}$

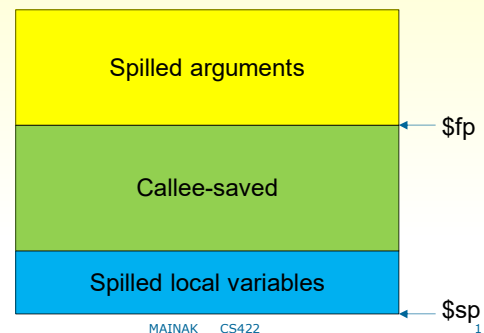
$\$ra$ — can be both callee or caller saved.
 deps on implementation.
 for main: it has to be callee saved before calling any other func.

Procedure call

- What does the caller do?
 - Save caller-saved registers if needed
 - Load arguments: first four in \$a0 to \$a3, rest on stack
 - jal or jalr
- What does the callee do? (will use frame pointer)
 - Save frame pointer at $-4(\$sp)$
 - Allocate stack for frame: $\$sp = \$sp - \text{frame size}$ (compiler knows the frame size for this procedure)
 - Save callee-saved registers if needed
 - Adjust frame pointer to point to the beginning of the frame: $\$fp = \$sp + \text{frame size} - 4$
- What happens on return?
 - Callee places return value in \$v0 (and \$v1 if 64-bit value)
 - Restore any callee-saved register from stack
 - Pop stack frame: $\$sp = \$sp + \text{frame size}$; jr \$ra

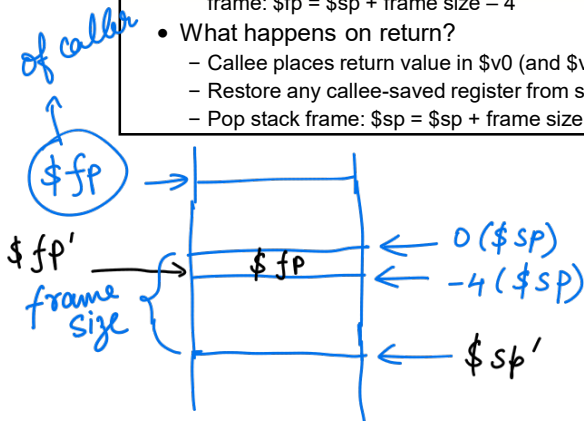
1

Procedure stack



MAINAK CS422

1



$\$sp = \text{TO S}$
 $\$fp = \text{BOS}$


```

main() {
    f() {
        g()
    }
}

```

caller

we don't
need reg \$ra

callee

main

↓

f

↓

g

f

↓

g

\$ra

\$ra m

\$ra +

\$ra m

\$ra f

mem

~~mem~~

~~mem~~

fact

by
main

AR

~~g AR~~

g AR